

Brian Lee

Code: <https://github.com/brialee/always-pizza/tree/data-storage-activity>

Method	Code Link	acs2015 time in seconds
Simple inserts	load_inserts.py	92.75
Drop Indexes and Constraints	load_inserts_without_index.py	90.27
Use UNLOGGED table	load_inserts_no_logging.py	13.43
Temp Table with memory tuning	load_inserts_temp_table.py	13.56
Batching	load_inserts_as_batch.py	10.31
copy_from	load_inserts_copy_from.py	1.41

J. Observations

Use this section to record any observations about the various methods/techniques that you used for bulk loading of the USA Census data. Did you learn anything about how and RDBMS functions and why various loading approaches produce varying performance results?

- Based on the outcomes above, it seems clear that the tools the DBMS vendor provides are the fastest under the given circumstances.
- I was a little surprised to see the reduction in runtime when removing indexing and constraints was as small as it was. I guess the alternative way to look at it is that adding those features doesn't add significant overhead for this dataset.
- An interesting finding was that disabling logging had comparable performance to using temporary tables with an increased memory buffer.